

O-Mem: Omni Memory System for Personalized, Long Horizon, Self-Evolving Agents

OPPO AI Agent Team

Abstract

Recent advancements in LLM-powered agents have demonstrated significant potential in generating human-like responses; however, they continue to face challenges in maintaining long-term interactions within complex environments, primarily due to limitations in contextual consistency and dynamic personalization. Existing memory systems often depend on semantic grouping prior to retrieval, which can overlook semantically irrelevant yet critical user information and introduce retrieval noise. In this report, we propose the initial design of O-Mem, a novel memory framework based on active user profiling that dynamically extracts and updates user characteristics and event records from their proactive interactions with agents. O-Mem supports hierarchical retrieval of persona attributes and topic-related context, enabling more adaptive and coherent personalized responses. O-Mem achieves 51.67% on the public LoCoMo benchmark, a nearly 3% improvement upon LangMem—the previous state-of-the-art, and it achieves 62.99% on PERSONAMEM, a 3.5% improvement upon A-Mem—the previous state-of-the-art. O-Mem also boosts token and interaction response time efficiency compared to previous memory frameworks. Our work opens up promising directions for developing efficient and human-like personalized AI assistants in the future.

Date: November 19, 2025

Correspondence: Wangchunshu Zhou at zhouwangchunshu@oppo.com

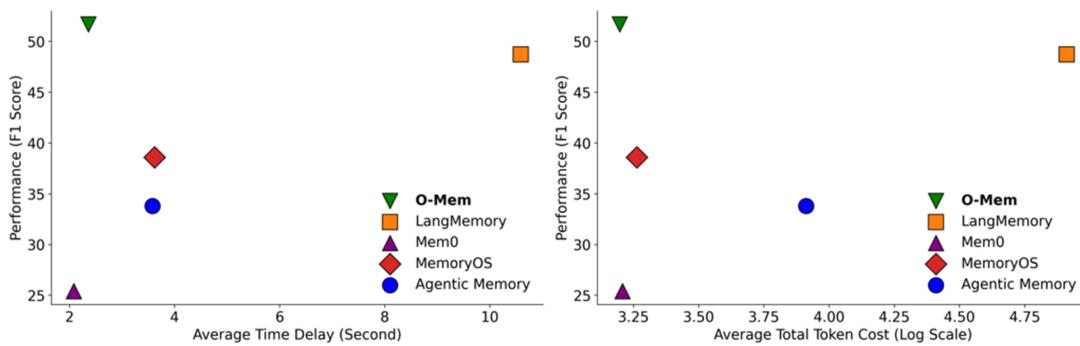


Figure 1 Trade-off between performance and efficiency of different memory frameworks. (a) Left panel: Average latency per interaction (MemoryOS latency uses FAISS-CPU for compatibility, a conservative estimate). (b) Right panel: Average computational cost (tokens) per interaction. Results show O-Mem achieves Pareto optimality in efficiency and performance. Note: Token-control experiments were only conducted on LoCoMo’s GPT-4.1; no token control for GPT-4o-mini and other two datasets.

1 Introduction

LLM-empowered agents have demonstrated huge potential in generating human-level intelligent responses [30] but still lack long-term interaction ability with complex external environments [41]. This limitation causes agents to struggle to maintain consistency of context across time [11] and reduced their personalization capability dynamically adapting to users' situations [40].

Agent memory systems equip agents with the ability to retain and utilize past experiences, unlike conventional agents that treat each interaction as independent. These systems store user past interactions in diverse architectures and enable agents to retrieve relevant information from them to deliver more personalized responses. For example, Memory OS [10] categorizes user interactions into short-term, mid-term, and long-term persona memory caches based on timestamps and frequency of occurrence. Agentic Memory [39] organizes interactions into distinct groups according to their semantic similarity, while Mem0 [5] extracts meaningful content from messages and stores the extracted information independently to support future retrieval. By structuring user information more effectively, these systems enhance the ability of agents to provide efficient and highly personalized responses.

The core pipeline of such memory systems involves grouping the messages based on semantic topics and retrieving memory groupings when interacting with users. However, this design presents several significant shortcomings: i) Memory systems that rely heavily on semantic retrieval may overlook information that is semantically irrelevant but potentially important—such as broader user characteristics or situational context—which is crucial for interactions requiring a comprehensive understanding of the user. As illustrated on the upper side of Figure 2, an intelligent agent should consider the user's health condition and recent schedule when planning weekend activities, rather than relying solely on activity-related memories. ii) The message grouping-based retrieval architecture can introduce additional retrieval noise. As shown on the lower side of Figure 2, sub-optimal memory groups can compel the agent to retrieve information from all three groups to gather sufficient context for appropriate responses. These redundant retrievals diminish the effectiveness of the model's responses, while also increasing latency and token consumption during LLM inference.

In this report, we propose the initial design of O-Mem, a human-centric memory framework based on active user profiling. Unlike conventional approaches that merely store and group past interactions between users and agents for retrieval, O-Mem actively extracts and updates user persona characteristics and event records from ongoing dialogs. This enables the agent to progressively refine its understanding of user attributes and historical experiences. We redefine personalized memory systems by treating each proactive user interaction as an opportunity for iterative user modeling. This approach effectively leverages both persona profiles and topic-related event records as contextual cues to support personalized responses. The main contributions of this work are as follows:

- We identify limitations in existing grouping-then-retrieval based semantic retrieval memory frameworks, notably their inadequate user understanding and restricted personalization abilities as they primarily depend on static historical interaction embeddings rather than constructing dynamic, multi-dimensional user contexts.
- We propose O-Mem, a novel persona memory framework that utilizes dynamic user profiling and a hierarchical, user-centric retrieval strategy. Unlike approaches that rely solely on semantic retrieval of past messages, O-Mem actively constructs and updates user profiles by accumulating knowledge from interaction histories.
- Extensive experiments on three persona-oriented tasks—i) persona-based open question answering, ii) persona-guided response selection and iii) persona-centric in-depth report generation—show that O-Mem consistently improves performance in a variety of memory applications. Specifically, O-Mem sets a new state-of-the-art in memory performance, achieving 51.76 % on the public LoCoMo benchmark and 62.99% on PERSONAMEM. By enabling dynamic user profiling and **interaction-time scaling**, O-Mem allows LLM agents to continuously adapt to users' evolving needs, demonstrating strong potential for enhancing long-term human-AI interactions through more personalized responses.

2 Related Work

Agent Memory System powered by large language models (LLMs) has gained significant popularity in recent years owing to their remarkable capabilities in task comprehension and execution [7, 15, 31, 38, 44, 45, 48]. Nevertheless, these systems continue to grapple with the challenge of sustaining high-quality performance when incorporating historical experience across complex, long-duration scenarios [10, 41]. To address this limitation, numerous agent memory

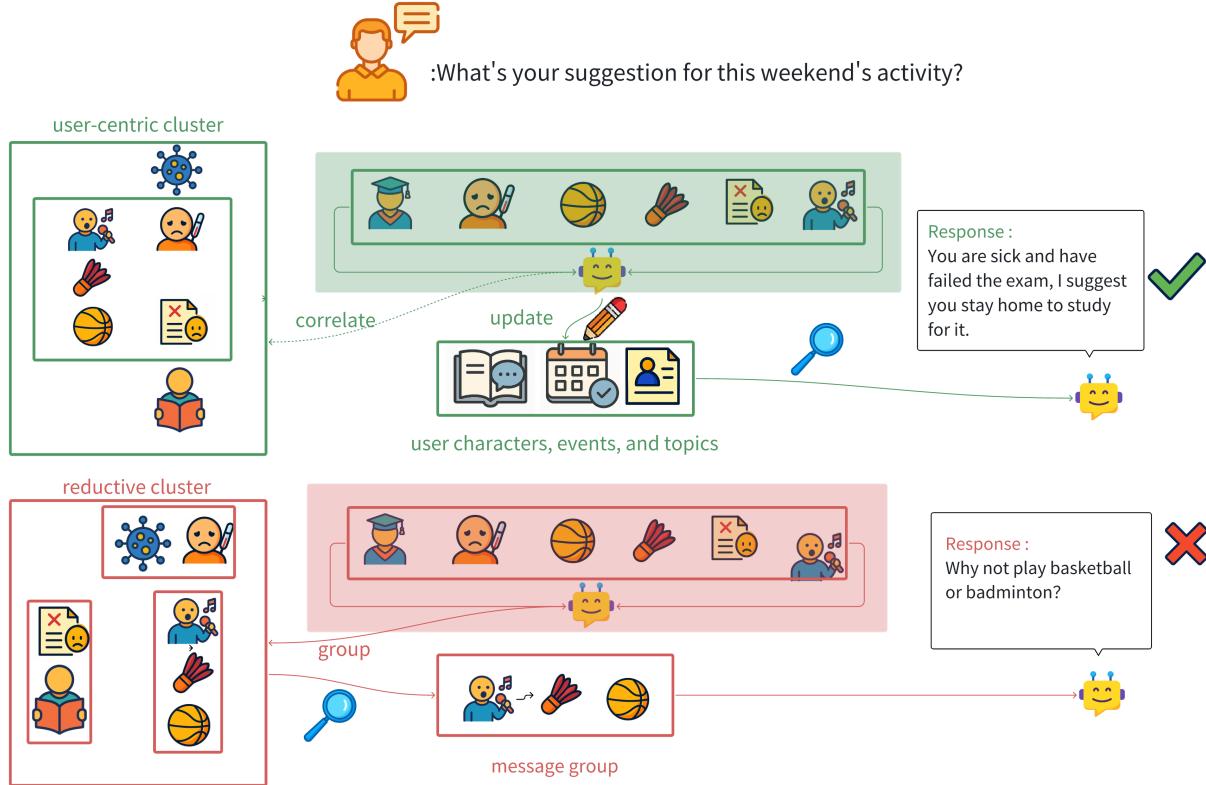


Figure 2 Top: our proposed user-centric framework O-Mem employing characteristic identification, event recording, and topic-message indexing. Bottom: the conventional memory system with semantic retrieval from message groupings. O-Mem correlates virtual relationships among user's interactions (in dotted line).

enhancement frameworks have been proposed, which can be broadly classified into two categories: (i) approaches that fine-tune LLM parameters to enhance information memorization and utilization [23, 37, 46], and (ii) methods that employ sophisticated information organization and retrieval techniques within external memory systems to preserve LLMs' long-term capabilities [43]. The latter approach has attracted considerable attention due to its plug-and-play nature, which eliminates the need for additional training costs. Furthermore, these methods significantly reduce the dependency of memory capacity on the LLMs' input window length. For example, Think-in-Memory (TiM) [20] preserves the reasoning traces of LLMs across multiple dialogue rounds to alleviate response inconsistencies. A-Mem [39] organizes memory fragments into linked lists to improve retrieval performance. Grounded Memory [24] introduces vision language models (VLMs) to interpret consecutive audio frames, organizing these interpretations in a graph structure for subsequent retrieval. MemoryBank [42] incorporates the Ebbinghaus Forgetting Curve theory to enable agents to forget and reinforce memories based on time elapsed and the relative significance of memory segments. MemGPT [25] and Memory OS [10] adopt an operating system-like architecture for memory organization and retrieval, employing mechanisms such as a first-in-first-out queue for working memory management. However, most existing systems overlook a critical aspect: how to dynamically and hierarchically establish connections between memory fragments to continuously update the agent's overall understanding of its environment. For instance, while A-Mem [39] and Memory OS [10] store semantically similar information in linked segments and retrieve the grouped data during response generation, their simple chunk-retrieval mechanisms, as illustrated in Figure 2, often fail to equip agents with a comprehensive and in-depth understanding of users prior to interaction. Therefore, to bridge this gap, we propose a novel memory system, O-Mem, based on active user profiling. The key difference between O-Mem and previous memory systems is that its core task is to answer the questions: "What kind of person is this user? What has he or she experienced?" rather than merely grouping received user information for later retrieval. We draw inspiration from the human brain's memory architecture and consequently redefine the three core components of an agent's memory: i). **Episodic Memory**, which is responsible for mapping a user's historical interaction cues to their corresponding situational

contexts (e.g., mapping the cue "project deadline" to the specific episode where the user expressed stress and requested help with scheduling); ii). **Persona Memory**, which constructs and maintains a holistic profile of the user; iii). **Working Memory**, which is responsible for providing relevant contextual information to the current interaction. Together, these components work in concert to enable O-Mem to build a deep, dynamic understanding of the user, powering more personalized and context-aware interactions. Figure 2 summarizes the key design features coherence of our approach in comparison with representative existing memory systems. We refer the conflict management as the process of memory systems to maintain the coherence of its stored information and agile user-centric user modeling as the process of timely iterating on the understanding of users.

Personalized Agent. While large language models (LLMs) serve as powerful assistants for a multitude of tasks, their effectiveness remains constrained without the ability to learn from and adapt to human preferences through personalization. A promising direction involves the development of persona agents—LLM-based systems deeply integrated with personal data to deliver responses aligned with user-specific needs [16]. Meeting the growing demand for such personalized interactions requires methodologies that can continuously and accurately infer user characteristics from their interactions [6, 21, 32, 36]. Most prior work on persona-enhanced LLMs has focused on injecting user information through fine-tuning [6, 29, 33, 47] or direct retrieval from user traces of static user profiles that rely on a limited set of predefined attributes[26, 28, 32, 34, 36]. However, these approaches face significant limitations in handling long-term, dynamic and evolving user preferences: fine-tuning requires computationally expensive retraining for each update, while direct retrieval lacks the capacity to synthesize longitudinal interaction patterns into coherent and evolving user profiles. In this work, we propose a persona memory system that dynamically organizes a user's interaction history into structured persona characteristics and experiential data, enabling more precise, adaptive, and personalized responses over time.

3 Method

When building the memory for O-Mem, we architect it to emulate a human-like memory model [12]. This is realized through three key properties: i).Long-Term Personality Modeling: It constructs a persistent and evolving user profile, mirroring the human ability to build a coherent understanding of others over time. ii).Dual-Context Awareness: It maintains both topical continuity (working memory) and associative, clue-triggered recall (episodic memory), enabling both coherent and precisely cued responses. iii).Structured, Multi-Stage Retrieval: It replaces a monolithic search with a structured process that orchestrates consultations with different memory types, resulting in more robust, transparent, and human-like reasoning. As illustrated in Figure 6, O-Mem continuously extracts and refines user profiles through ongoing interaction, building a semantic mapping between topics/clues and corresponding interaction scenarios. This enables dynamic, multi-faceted user understanding that supports powerful personalization. In this section, we first present the notation used in our work and the storage formats of different memory components (subsection 3.1), followed by an explanation of how user interactions are encoded into the different memory components in O-Mem (subsection 3.2), and finally describe the retrieval process across these memories (subsection 3.3).

3.1 Preliminary: Notation and Memory Components Storage Format

In this section, we define a user interaction, denoted as U , as a record of either explicit literal content (e.g., search queries) or implicit user behavior (e.g., taking a screenshot). Let M_w be a dictionary that maps each clue word w to the interactions in which it appears, and M_t be a dictionary that maps each topic to its corresponding interactions. Additionally, P_a denotes the list of persona attributes, and P_f represents the list of persona fact events:

Let $\mathcal{U} = \{U_1, U_2, \dots, U_n\}$ be the set of user interactions.

$$M_w[w] = \{U \in \mathcal{U} \mid w \text{ appears in } U\}, \quad M_t[t] = \{U \in \mathcal{U} \mid U \text{ is associated with topic } t\}.$$

As illustrated in Figure 6, we model these components within a cognitive architecture: P_f and P_a constitute the **user persona memory**, which stores long-term, abstracted user knowledge; the **topic-interaction dictionary** M_t functions as **working memory**, capturing the topical context of the current interaction; and the **keyword-interaction dictionary** M_w serves as **episodic memory**, acting as an associative index that links salient clues to their originating interactions. Unlike the strict physiological definitions of working memory and episodic memory, **The definitions of agent working memory and episodic memory in O-Mem are past interactions related to the current interaction topic and past interactions**

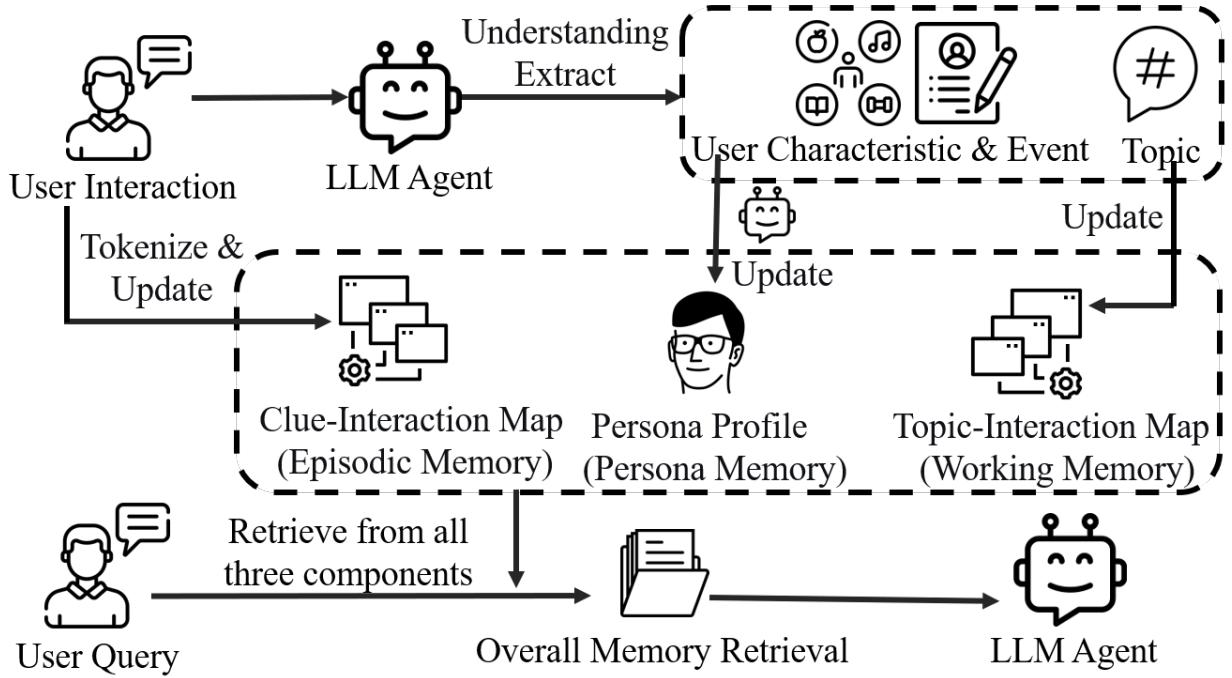


Figure 3 Top: The process of encoding user interactions into memory in O-Mem. Different colors refers to different memory components. O-Mem encodes a user interaction into memory by extracting and recording relevant user attributes and event data into persona memory,episodic memory, and working memory. Bottom: The memory retrieval process concerning one user interaction in O-Mem. O-Mem retrieves from all its three memory components concerning one new user query.

related to clues in the current interaction, respectively. The semantic similarity function $s(t_1, t_2)$ between two text segments t_1 and t_2 , and the memory retrieval function $F_{\text{Retrieval}}$ based on s , are formally defined as follows:

$$s(t_1, t_2) = \frac{\mathbf{f}_e(t_1) \cdot \mathbf{f}_e(t_2)}{\|\mathbf{f}_e(t_1)\| \|\mathbf{f}_e(t_2)\|}, \quad F_{\text{Retrieval}}(M | q) = \text{top-}k\{s(m, q) | m \in M\}$$

where $f_e(\cdot)$ denotes a text embedding function, top- k returns the k most similar items, and M refers to the memory component from which retrieval is performed.

3.2 Memory Construction Process

There are three components of O-Mem: i).**Persona Memory:** This component stores the user’s long-term attributes and significant factual events in a structured profile. Its function is to enable Long-Term Personality Modeling, ensuring responses are consistently personalized by maintaining a persistent understanding of the user’s identity. ii).**Working Memory:** This component maintains a dynamic index mapping conversation topics to related interactions. Its function is to support dual-context awareness by preserving topical continuity, thereby ensuring dialogue coherence through relevant contextual grounding. iii).**Episodic Memory:** This component serves as an associative index linking specific clues or keywords to their original interactions. Its function is to enable associative, clue-triggered recall as part of dual-context awareness, allowing precise retrieval of specific events beyond mere semantic similarity. The three memory components employ distinct technical approaches tailored to their specific functions: persona memory uses LLM-based extraction to detect user attributes and events, with regulation through a decision process (Add/Ignore/Update) to maintain profile coherence. Working memory automatically indexes interactions under LLM-identified topics, relying primarily on accurate topic detection for regulation. Episodic memory tokenizes interactions to detect potential clues, regulated by a distinctiveness filter that prioritizes rare keywords. This statistics-driven approach enables precise associative recall. The differentiation in update mechanisms stems from their distinct purposes: persona memory requires careful curation, while working and episodic memories benefit from automated indexing for efficiency.

Given the i -th user interaction u_i , O-Mem first extracts its topic t_i , revealed user attribute a_i and past event e_i with language model \mathcal{L} :

$$(t_i, a_i, e_i) = \mathcal{L}(u_i) \quad (1)$$

The clue interaction map M_w and topic interaction map M_t are updated by increasing the count for each word in u_i and t_i :

$$M_t^{(i+1)}[t_i] \leftarrow M_t^{(i)}[t_i] \cup \{i\}, \quad M_w^{(i+1)}[w_j] \leftarrow M_w^{(i)}[w_j] \cup \{i\}, \quad \forall w_j \in \mathcal{T}(u_i) \quad (2)$$

where $\mathcal{T}(u_i) = \{w_1, w_2, \dots, w_n\}$ represents the tokenized words from the i -th user interaction. For e_i , \mathcal{L} generates an memory management operation regarding its integration with the existing persona fact event list P_f :

$$\text{Op}(e_i) \leftarrow \mathcal{L}(e_i, P_f) \in \{\text{Add}, \text{Ignore}, \text{Update}\}, \quad P_f \leftarrow \text{ApplyOp}(P_f, e_i, \text{Op}(e_i)) \quad (3)$$

where Op refers to the operation decision from \mathcal{L} and ApplyOp refers to the function that executes this operation. During our observation, we identified that similar attributes from the same user frequently recur across different interactions (e.g., users always mention their hobbies repeatedly). To better organize these extracted attributes, we propose an LLM-augmented nearest neighbor clustering method:

$$\text{Op}(a_i) \leftarrow \mathcal{L}(a_i, P_a^t) \in \{\text{Add}, \text{Ignore}, \text{Update}\}, \quad P_a^t \leftarrow \text{ApplyOp}(P_a^t, a_i, \text{Op}(a_i)) \quad (4)$$

$$\text{NN}(a_i) = \arg \min_{a_l \in P_a^t, l \neq i} (1 - s(a_i, a_l)) \quad (5)$$

$$G = (V, E), \quad V = \{a_1, \dots, a_K\}, \quad E = \{(a_l, \text{NN}(a_l)) \mid a_l \in P_a^t\} \quad (6)$$

$$\mathcal{B} = \{B_1, \dots, B_M\} = \text{ConnectedComponents}(G), \quad P_a = \bigcup_{m=1}^M \mathcal{L}(B_m) \quad (7)$$

where $\text{NN}(a_i)$ denotes the nearest neighbor of attribute a_i based on the similarity function $s(\cdot, \cdot)$; $G = (V, E)$ represents the nearest-neighbor graph constructed from the temporary attribute list P_a^t , with vertices corresponding to attributes and edges connecting each attribute to its nearest neighbor; K refers to the total number of attributes in P_a^t , $\mathcal{B} = B_1, \dots, B_M$ is the set of connected components obtained from G via connected component analysis; and the final attribute set P_a is obtained by applying the large language model \mathcal{L} to analyze the aggregated attributes within each connected component.

3.3 Memory Retrieval Process

O-Mem employs a parallel retrieval strategy across all three memory components. The system simultaneously queries the Persona Memory for user profile context, Working Memory for topical context, and Episodic Memory for clue-triggered events. The retrieved information from all memory components is then aggregated and processed by the language model to generate the final response. This parallel approach ensures comprehensive context integration while maintaining the distinct advantages of each component. For each user interaction u_i , O-Mem conducts retrieval from the user's persona memory, episodic memory, and working memory. We introduce their retrieval process separately.

Working Memory Retrieval. we define the retrieval process of working memory as:

$$R_{\text{working}} = \bigcup_{t \in \hat{T}} M_t[t], \quad \text{where } \hat{T} = F_{\text{Retrieve}}(\mathcal{K}(M_t), u_i) \quad (8)$$

where M_t is the mapping from topics to their corresponding interactions, $\mathcal{K}(M_t)$ denotes the set of topics in M_t , u_i is the current interaction, F_{Retrieve} retrieves the most relevant topics \hat{T} for u_i from $\mathcal{K}(M_t)$, and R_{working} is the set of interactions in M_t corresponding to these relevant topics.

Episodic Memory Retrieval. We define the retrieval process of episodic memory as follows. The episodic memory is structured as a word-to-interactions mapping dictionary M_w ($M_w : w \rightarrow \{i\}$), which maps words to the sets of past interactions (memory entries) in which they appear. That is, for a word w , $M_w[w]$ yields all past interactions containing w . Given the current user interaction u_i , the retrieval process is: (1) **Tokenization:** Tokenize the utterance into a sequence of words: $W = \text{Tokenize}(u_i)$. (2) **Clue Selection:** Calculate the clue selection score for each word $w \in W$ with respect to the clue-interaction map M_w . The word with the highest score is selected as the target clue \hat{w} :

$$\hat{w} = \arg \max_{w \in W} \text{Score}(w, M_w) \quad (9)$$

$$\text{Score}(w, M_w) = \frac{1}{d_f w} \quad (10)$$

where df_w is the number of past interactions in M_w that contain the word w (i.e., $df_w = |M_w[w]|$). The set of episodic memory interactions associated with the clue \hat{w} is then retrieved as: $R_{episodic} = M_w[\hat{w}]$.

Persona Memory Retrieval. We define the persona retrieval process as:

$$R_{persona} = F_{\text{Retrieval}}(P_f, u_i) \oplus F_{\text{Retrieval}}(P_a, u_i) \quad (11)$$

where P_f refers to the persona facts, P_a refers to the persona attributes, u_i is the current user interaction, \oplus denotes the concatenation operation, and $R_{persona}$ refers to the retrieved persona information.

Overall Memory Retrieval. We define the overall retrieval and final response as:

$$R = R_{\text{working}} \oplus R_{\text{episodic}} \oplus R_{\text{persona}}, O = \mathcal{L}(R, u_i) \quad (12)$$

where R represents the overall retrieved memory content, O represents the final response generated by the language model \mathcal{L} based on the current user interaction u_i and the retrieved memories.

4 Experiment

Datasets and Evaluation Metrics. We evaluate our method on three benchmarks: **LoCoMo** [22], **PERSONAMEM** [9], and **Personalized Deep Research Bench** [18].

The **LoCoMo** benchmark features extended dialogues averaging 300 turns across four memory challenge types: Single-hop, Multi-hop, Temporal, and Open-domain. The **PERSONAMEM** dataset contains user-LLM conversations spanning 15 diverse topics. To address the need for evaluating personalized long-text generation, we introduce **Personalized Deep Research Bench**, a benchmark simulating real-world deep research scenarios [18]. Unlike existing datasets, Personalized Deep Research Bench comprises 50 deep research queries derived from multi-round conversations between 25 real users and LLMs, requiring nuanced understanding of individual user characteristics. It is built upon a subset of a persona deep research dataset collected from real users through commercial applications but is specifically repurposed and curated by the dataset construction committee for assessing memory system¹. For evaluation, we employ:

- **LoCoMo**: F1 and BLEU-1 scores following the standard protocol;
- **PERSONAMEM**: Accuracy for multiple-choice questions;
- **Personalized Deep Research Bench**: Goal Alignment and Content Alignment scores, measuring adherence to user characteristics and expectations via LLM-as-a-judge.

Compared Baselines. Our method is compared with: (i) open-source memory frameworks: A-Mem [39], MemoryOS [10], Mem0 [5], and LangMem [13]; and (ii) commercial/proprietary frameworks: ZEP [27], Memos [17], and OpenAI [1]. Due to budget constraints and licensing costs, we report results from original publications for commercial frameworks. Mem0 was evaluated using its open-source version due to cost and accessibility.

Implementation Details. We use all-MiniLM-L6-v2 [2] as embedding model in O-Mem to calculate similarities. All of our experiments are conducted on two A800 GPUs. The choice of language models across datasets was informed by computational budget. A comparative analysis using both GPT-4.1 and GPT-4o-mini was performed on the LoCoMo benchmark. For the remaining datasets (PERSONAMEM and Personalized Deep Research Bench), only the larger GPT-4.1 model was used. Due to resource constraints, the experiments were conducted on public shared servers. The hardware resources in cloud environments (e.g., GPU computational capacity, network latency) may exhibit inherent volatility, meaning we cannot guarantee identical latency and computational states across experimental runs. The generative nature of Large Language Models (LLMs) is inherently stochastic. Given the significant economic cost associated with API calls, it was not feasible to perform multiple repeated experiments to obtain statistical summaries (e.g., mean and standard deviation). Furthermore, in pursuit of evaluation fairness, we deliberately avoided fixing the random seed to prevent any potential—even unintentional—"seed cherry-picking," thereby ensuring the reported result represents a single sample from the distribution of possible outcomes. Therefore, we emphasize that the primary

¹The full original Personalized Deep Research Bench benchmark has been released to the community comprises the entire initial query set, including the 50 highly personalized memory-related queries selected for this study as well as the broader collection, which, despite being less discriminative for memory personalization, remains a valuable asset for future research. Specifically, the data construction committee removed some queries that were too difficult or too easy for the memory system to answer from the original deep research dataset, making it more cost-effective.

Table 1 Performance comparison using different LLMs on LoCoMo with best scores highlighted.

| LLM | Method | Cat1: Multi-hop | | Cat2: Temporal | | Cat3: Open | | Cat4: Single-hop | | Average | |
|-------------|----------|-----------------|--------------|----------------|--------------|--------------|--------------|------------------|--------------|--------------|--------------|
| | | F1 | B1 | F1 | B1 | F1 | B1 | F1 | B1 | F1 | B1 |
| GPT-4.1 | LangMem | 41.11 | 32.09 | 53.67 | 46.22 | 33.38 | 27.26 | 51.13 | 44.22 | 48.72 | 41.36 |
| | Mem0 | 30.45 | 22.15 | 10.69 | 9.21 | 16.75 | 11.34 | 30.32 | 25.82 | 25.40 | 20.78 |
| | MemoryOS | 29.25 | 20.79 | 37.73 | 33.17 | 22.70 | 18.65 | 43.85 | 38.72 | 38.58 | 33.03 |
| | A-Mem | 29.29 | 21.47 | 33.12 | 28.50 | 15.41 | 12.34 | 37.64 | 32.88 | 33.78 | 28.60 |
| | Ours | 42.64 | 34.08 | 57.48 | 49.76 | 30.58 | 25.69 | 54.89 | 48.98 | 51.67 | 44.96 |
| GPT-4o-mini | LangMem | 36.03 | 27.22 | 38.10 | 32.23 | 29.79 | 23.17 | 41.72 | 35.61 | 39.18 | 32.59 |
| | Mem0 | 17.19 | 12.06 | 3.59 | 3.37 | 12.24 | 8.57 | 12.74 | 10.62 | 11.62 | 9.24 |
| | ZEP | 23.14 | 14.96 | 17.59 | 14.57 | 19.76 | 13.17 | 32.49 | 27.38 | 26.88 | 21.55 |
| | MemoryOS | 41.15 | 30.76 | 20.02 | 16.52 | 48.62 | 42.99 | 35.27 | 25.22 | 34.00 | 25.53 |
| | OpenAI | 33.10 | 23.84 | 23.90 | 18.25 | 17.19 | 11.04 | 36.96 | 30.72 | 32.30 | 25.63 |
| | A-Mem | 33.23 | 29.11 | 8.04 | 7.81 | 34.13 | 27.73 | 22.61 | 15.25 | 22.24 | 17.02 |
| | MEMOS | 35.57 | 26.71 | 53.67 | 46.37 | 29.64 | 22.40 | 45.55 | 38.32 | 44.42 | 36.88 |
| Ours | | 44.17 | 34.78 | 53.54 | 45.65 | 25.24 | 19.22 | 54.53 | 48.33 | 50.60 | 43.48 |

Table 2 Performance comparison on PERSONAMEM with GPT-4.1.

| Method | Recall user shared facts | Suggest new ideas | Track full preference evolution | Revisit reasons behind preference updates | Provide preference-aligned recommendations | Generalize to new scenarios | Average |
|-----------|--------------------------|-------------------|---------------------------------|---|--|-----------------------------|--------------|
| LangMem | 31.29 | 24.73 | 53.24 | 81.82 | 40.00 | 8.77 | 42.61 |
| Mem0 | 32.13 | 15.05 | 54.68 | 80.81 | 52.73 | 57.89 | 46.86 |
| A-Mem | 63.01 | 27.96 | 54.68 | 85.86 | 69.09 | 57.89 | 59.42 |
| Memory OS | 72.72 | 17.20 | 58.27 | 78.79 | 72.72 | 56.14 | 58.74 |
| O-Mem | 67.81 | 21.51 | 61.15 | 89.90 | 65.45 | 73.68 | 62.99 |

contribution of this work is to demonstrate the feasibility and fundamental trends of the proposed method, rather than to provide a highly precise performance benchmark under controlled conditions. We encourage readers to focus their evaluation on the relative trends within the same unstable environment and to interpret the absolute performance values with caution. Owing to regional constraints, access to the GPT model was facilitated through an intermediary API provider. However, the reliability of this service proved erratic, manifesting as frequent request failures and malformed responses. These issues, which were exacerbated during large-scale data construction, introduced significant uncertainty into our estimates of computational cost and time expenditure. Consequently, this variability precluded us from reporting precise values for these metrics. Mem0 was evaluated using its open-source version due to cost and accessibility. For efficiency, the GPT-4o results for baselines (excluding Mem0) are adopted from existing literature.

Table 3 Performance comparison on Personalized Deep Research Bench with GPT-4.1.

| Method | Goal Alignment | Content Alignment | Average |
|-----------|----------------|-------------------|--------------|
| | Alignment | Alignment | |
| Mem0 | 37.32 | 35.54 | 36.43 |
| Memory OS | 40.60 | 39.67 | 40.14 |
| O-Mem | 44.69 | 44.29 | 44.49 |

Performance Comparison. The experimental results in three benchmark datasets are separately shown in [Table 1](#), [Table 2](#), and [Table 3](#). Due to limited access to ZEP[27], Memos[17], and OpenAI memory[1], we only report their performance reported in their work using GPT-4o-mini. For Personalized Deep Research Bench benchmark dataset, we only compare our method with mem0[5] and MemoryOS[10] for cost efficiency. O-Mem demonstrates superior performance compared to all baselines across three benchmark datasets. The performance advantage is more pronounced in complex reasoning tasks. As shown in [Table 1](#), on the comprehensive LoCoMo benchmark, O-Mem achieves the highest average F1 scores of 51.67% with GPT-4.1 and 50.60% with GPT-4o-mini, outperforming the strongest baselines by significant margins (2.95% and 6.18% absolute improvements, respectively). The performance

Table 4 Performance and efficiency comparison between direct retrieval from complete raw interaction history (Direct RAG) and O-Mem.

| Method | F1 (%) | Avg. Token Cost | Peak Memory Overhead (MB) ³ | Delay (s) |
|------------|--------|-----------------|--|-----------|
| Direct RAG | 50.25 | 2.6K | 33.16 | 4.01 |
| O-Mem | 51.67 | 1.5K | 22.99 | 2.36 |

*For a fair comparison, the reported overhead for both RAG and O-Mem is calculated as the peak GPU memory usage minus the fixed memory allocated by the same embedding model used in our main experiments.

advantage is particularly pronounced in complex reasoning tasks. For Temporal reasoning, O-Mem achieves F1 scores of 57.48% (GPT-4.1) and 53.54% (GPT-4o-mini), substantially outperforming all baselines. This indicates that our memory management mechanism effectively handles temporal dependencies and sequential information, which is crucial for maintaining coherent long-term conversations. Table 2 further demonstrates O-Mem’s effectiveness in personalized interaction scenarios on the PERSONAMEM dataset. O-Mem achieves an average accuracy of 62.99%, exceeding the closest competitor (A-Mem at 59.42%) by 3.57%. Notably, O-Mem excels in challenging tasks such as "Generalize to new scenarios" (73.68%) and "Revisit reasons behind preference updates" (89.90%), highlighting its robust capability in understanding and adapting to evolving user preferences. The superiority of O-Mem is consistently validated on our newly introduced Personalized Deep Research Bench dataset (Table 3), where it achieves an average alignment score of 44.49%, significantly higher than Mem0 (36.43%). This 8.06% improvement demonstrates our method’s practical utility in real-world personalized deep research scenarios that require nuanced understanding of user characteristics. A fair comparison was conducted by generating all deep research reports through the centralized sonar-deep-research service [3], leveraging retrievals from each method’s individual memory system.

5 Discussion

Rethinking the Value of Memory Systems. Do we truly need meticulously designed, complex memory systems? Most existing approaches adhere to a common paradigm: during retrieval, systems access processed user interactions rather than raw historical data. This design is largely driven by increasingly stringent privacy regulations worldwide [4, 8, 35]. By relying on abstracted user data, memory systems help AI companies **mitigate legal risks** while maintaining personalization capabilities. However, this abstraction comes at a significant cost: **the irreversible loss of information fidelity and contextual nuance.** For instance, a detailed user statement such as "*admiring a specific lamp and rug in a downtown antique store last Saturday*" may be compressed into a structured preference like "*[User] likes vintage home decor.*" While efficient, this compression sacrifices granular details—specific objects, locations, and temporal context—that are crucial for precise and contextually relevant interactions.

To quantify this trade-off, we compare the performance of direct retrieval-augmented generation (RAG) overall raw interactions² with O-Mem. As shown in Table 4, direct RAG achieves competitive performance (50.25 vs. 51.67 F1 score) despite its conceptual simplicity, though at a substantially higher computational cost (2.6K vs. 1.5K tokens). Notably, when compared to the results in Table 1, direct RAG with access to complete interaction history achieves competitive performance, highlighting the fundamental value of preserving raw interaction data. However, this comes at prohibitive computational costs that limit practical deployment. O-Mem addresses this critical limitation by achieving comparable performance with significantly reduced overhead, positioning it as a computationally efficient alternative that balances performance with practicality.

As depicted in Table 4, we also evaluate the practical deployment advantages of O-Mem by measuring average peak retrieval memory overhead [19] per response. O-Mem achieves a significant 30.6% reduction in peak memory overhead (from 33.16 MB to 22.99 MB), substantially relaxing hardware constraints for large-scale personalized inference. While direct RAG processes complete interaction histories verbatim, O-Mem maintains distilled representations that preserve semantic essence while drastically reducing sequence length. This design choice yields substantial computational memory benefits. For response latency, O-Mem demonstrates a 41.1% reduction in delay compared to direct RAG (from 4.01 seconds to 2.36 seconds), highlighting its efficiency for real-time applications.

²This includes all original interactions. To avoid missing contextual information in direct retrieval, the responses from the agent are also retrieved, which is different from our work that focuses on modeling user interactions.

³*

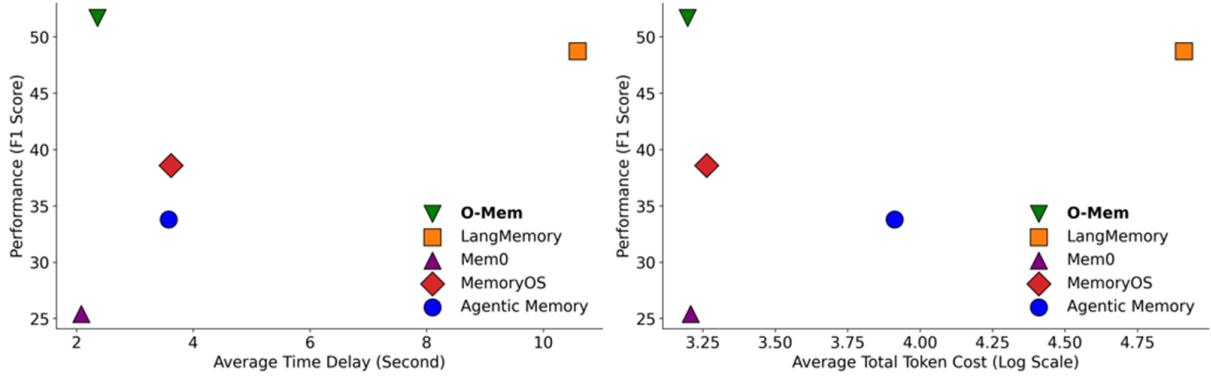


Figure 4 Trade-off between performance and efficiency of different memory frameworks. The left panel (a) compares the average latency per interaction. The MemoryOS latency was evaluated using FAISS-CPU due to compatibility issues on our computing platforms, thus representing a conservative estimate of its latency. (b) The right panel compares the average computational cost (in tokens) per interaction. Results demonstrate that O-Mem achieves a Pareto-optimal solution in both efficiency and overall performance. Note that we only performed token-control experiments on Locomo’s GPT-4.1 experiments. We did not control tokens for the experiments on GPT-4o-mini and the other two datasets.

Efficiency Analysis. We evaluate the efficiency of O-Mem by measuring the average token consumption and latency per response on the LoCoMo benchmark. The results, presented in Figure 4, substantiate that O-Mem achieves a superior balance between efficiency and effectiveness. Compared to the highest-performing baseline, LangMem (48.72 F1), O-Mem (51.67 F1) reduces token consumption by **94%** (from 80K to 1.5K tokens) and latency by **80%** (from 10.8s to 2.4s), while delivering superior performance. Against the **second-best performing** baseline, MemoryOS (38.58 F1), O-Mem not only secures a **34%** higher F1 score but also reduces latency by **34%** (2.4s vs. 3.6s). These results unequivocally demonstrate that O-Mem sets a new Pareto frontier for efficient and effective memory systems.

The efficiency advantage of O-Mem stems from two key design choices: The first is the independence of its retrieval operations across the three memory components. Unlike sequential architectures (e.g., A-men) that rely on a cascade of coarse-to-fine stages, O-Mem performs a one-time, concurrent retrieval across all three memory paths. Secondly, the retrieval mechanism of O-Mem utilizes user persona information, which, as opposed to raw user interaction records, typically contains less noise, thereby enhancing the cost-effectiveness of token usage. O-Mem also achieves a substantial reduction in **storage footprint**, requiring only nearly **3 MB per user**—much less than the nearly **30 MB per user** consumed by Memory OS. This storage efficiency stems from our topic/keyword-based mapping design, which utilizes a lightweight index, in contrast to Memory OS which must store dense vector mappings for each memory chunk. Furthermore, O-Mem employs a radically simplified inference pipeline. Each response is generated through **only one LLM invocation** (three times for LangMemory). This streamlined workflow enables O-Mem to achieve superior efficiency with minimal reponse latency and computational expense.

Memory Component Analysis To quantify the contribution of each core module in our framework, we performed an ablation study on all three memory components—Persona Memory (PM), Episodic Memory (EM), and Working Memory (WM)—using the LoCoMo benchmark. The results are summarized in Table 5. As indicated in the first three rows of the table, each module individually contributes to improved overall performance. However, **such performance gains could be partially attributed to the increased volume of retrieved information**, which leads to longer retrieval sequences and higher token consumption during response generation. This trade-off between performance and efficiency has often been overlooked in prior ablation studies of memory-augmented systems. To definitively isolate this confound, we conducted a token-controlled ablation study (Rows 4–5 in Table 5), wherein the total token budget for each ablated configuration was fixed at 1.5K tokens to match that of the full O-Mem framework (WM+EM+PM). The clear performance gradient under a fixed token budget provides conclusive evidence that the performance gains are attributable to the *quality and relevance* of the information retrieved by each module, not merely to an increase in context. This finding confirms that each independent memory module of O-Mem effectively captures distinct and complementary aspects of the interactions.

Table 5 Ablation study on different components of O-Mem using the LoCoMo benchmark dataset.

| Memory Configuration | F1 (%) | Bleu-1 (%) | Total Tokens |
|----------------------------|--------|------------|--------------|
| WM only | 44.03 | 38.05 | 1.3K |
| WM + EM | 49.62 | 43.18 | 1.4K |
| WM + EM + PM | 51.67 | 44.96 | 1.5K |
| WM + EM (token-controlled) | 50.10 | 43.27 | 1.5K |
| WM only (token-controlled) | 46.07 | 39.95 | 1.5K |

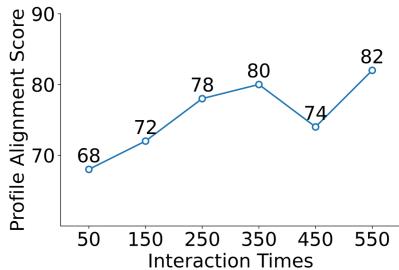


Figure 5 Memory profile alignment dynamics during memory-time Scaling. More interactions lead to more concise user understanding from O-Mem.

| Memory Configuration | Average Performance | Average Retrieval Length (Chars) |
|----------------------|---------------------|----------------------------------|
| O-Mem | 44.49 | 6499 |
| O-Mem w/o Attributes | 42.14 | 28555 |

Table 6 Ablation study on the impact of persona attributes, evaluated on the Personalized Deep Research Bench. Removing attributes from O-Mem not only causes a performance drop but, more notably, leads to a substantial increase in retrieval length, indicating that attributes are crucial for precise memory filtering. This demonstrates the effectiveness and efficiency gains brought by user attribute extraction.

Memory-Time Scaling for User Understanding. We conduct a systematic evaluation of O-Mem’s user understanding capability by examining how it **scales with the number of interactions** through two key analyses: (1) verifying the accuracy of persona attributes extracted from interaction data, and (2) assessing the practical utility of these attributes in personalizing the agent’s responses. First, to evaluate the **scaling of extraction accuracy**, we collect persona attributes inferred by O-Mem from a single user’s dialogue history **across increasing interaction counts**. These extracted attributes are then compared against the user’s ground-truth profile using an LLM-as-judge scoring mechanism [14] to measure alignment. To evaluate the fidelity of the user profiles extracted by O-Mem, we have designed structured prompts that direct a large language model to compare the extracted profiles against the ground-truth user profiles and assign a consistency score. We acknowledge the inherent limitations of this method in achieving perfect objectivity. Its primary objective, however, is not to provide a precise measurement but to fundamentally assess whether O-Mem can construct a dynamic and increasingly accurate user portrait as the number of interactions scales. As shown in Figure 5, the extracted persona attributes gradually converge toward the ground-truth profile **as interactions scale**, demonstrating that O-Mem effectively refines its user understanding **through this scaling process**. Second, to measure the practical impact of persona attributes, we compare O-Mem with and without access to persona attributes on the Personalized Deep Research Bench dataset. Results in Table 6 show that incorporating persona attributes yields a significant improvement in response personalization (average performance increases from 42.14 to 44.49) while substantially reducing the retrieval length (from 28,555 to 6,499 characters). These results demonstrate that O-Mem’s ability to extract and leverage user attributes **through scaled interactions** substantially enhances its performance in complex personalized text generation tasks, achieving stronger personalization with improved efficiency.

6 Conclusion

In this paper, we propose O-Mem, a novel memory framework that enhances long-term human-AI interaction through dynamic user profiling and hierarchical memory retrieval. Unlike conventional approaches that rely solely on semantic retrieval of past messages, O-Mem actively constructs and refines user profiles from ongoing interactions. This approach effectively addresses the key limitations of conventional methods in maintaining long-term, consistent user context. Extensive experiments on three personalized benchmarks demonstrate that O-Mem achieves state-of-the-art performance while reducing token consumption by 94% and inference latency by 80% compared to its closest competitor, highlighting its superior efficiency. The proposed framework provides an effective solution for complex personalized text generation tasks, enabling LLM agents to deliver more coherent and contextually appropriate responses. Our work opens up promising directions for developing more efficient and human-like personalized AI assistants in the future.

References

- [1] URL <https://openai.github.io/openai-agents-python/ref/memory/>.
- [2] 2024. URL <https://huggingface.co/sentence-transformers/all-MiniLM-L6-v2>.
- [3] 2025. URL <https://docs.perplexity.ai/getting-started/models/models/sonar-deep-research>.
- [4] Igor Calzada. Citizens' data privacy in china: The state of the art of the personal information protection law (pipl). *Smart Cities*, 5(3):1129–1150, 2022.
- [5] Prateek Chhikara, Dev Khant, Saket Aryan, Taranjeet Singh, and Deshraj Yadav. Mem0: Building production-ready ai agents with scalable long-term memory. *arXiv preprint arXiv:2504.19413*, 2025.
- [6] Joel Eapen and VS Adhithyan. Personalization and customization of llm responses. *International Journal of Research Publication and Reviews*, 4(12):2617–2627, 2023.
- [7] Mehmet Firat and Saniye Kuleli. What if gpt4 became autonomous: The auto-gpt project and use cases. *Journal of Emerging Computer Technologies*, 3(1):1–6, 2023.
- [8] Henry Hosseini, Christine Utz, Martin Degeling, and Thomas Hupperich. A bilingual longitudinal analysis of privacy policies measuring the impacts of the gdpr and the ccpa/cpra. 2024.
- [9] Bowen Jiang, Zhuoqun Hao, Young-Min Cho, Bryan Li, Yuan Yuan, Sihao Chen, Lyle Ungar, Camillo J Taylor, and Dan Roth. Know me, respond to me: Benchmarking llms for dynamic user profiling and personalized responses at scale. *arXiv preprint arXiv:2504.14225*, 2025.
- [10] Jiazheng Kang, Mingming Ji, Zhe Zhao, and Ting Bai. Memory os of ai agent. *arXiv preprint arXiv:2506.06326*, 2025.
- [11] Philippe Laban, Hiroaki Hayashi, Yingbo Zhou, and Jennifer Neville. Llms get lost in multi-turn conversation. *arXiv preprint arXiv:2505.06120*, 2025.
- [12] Runchen Lai, Youjun Li, Fan Yang, Ying Li, Nan Yao, Chunwang Su, Si-Ping Zhang, Yuanyuan Mi, Celso Grebogi, and Zi-Gang Huang. A brain-inspired neurodynamic model for efficient sequence memory. *Neurocomputing*, page 131849, 2025.
- [13] Langchain-Ai. Github - langchain-ai/langmem. URL <https://github.com/langchain-ai/langmem>.
- [14] Haitao Li, Qian Dong, Junjie Chen, Huixue Su, Yujia Zhou, Qingyao Ai, Ziyi Ye, and Yiqun Liu. Llms-as-judges: a comprehensive survey on llm-based evaluation methods. *arXiv preprint arXiv:2412.05579*, 2024.
- [15] Weizhen Li, Jianbo Lin, Zhusong Jiang, Jingyi Cao, Xinpeng Liu, Jiayu Zhang, Zhenqiang Huang, Qianben Chen, Weichen Sun, Qixiang Wang, Hongxuan Lu, Tianrui Qin, Chenghao Zhu, Yi Yao, Shuying Fan, Xiaowan Li, Tiannan Wang, Pai Liu, King Zhu, He Zhu, Dingfeng Shi, Piaohong Wang, Yeyi Guan, Xiangru Tang, Minghao Liu, Yuchen Eleanor Jiang, Jian Yang, Jiaheng Liu, Ge Zhang, and Wangchunshu Zhou. Chain-of-agents: End-to-end agent foundation models via multi-agent distillation and agentic rl, 2025. URL <https://arxiv.org/abs/2508.13167>.
- [16] Yuanchun Li, Hao Wen, Weijun Wang, Xiangyu Li, Yizhen Yuan, Guohong Liu, Jiacheng Liu, Wenxing Xu, Xiang Wang, Yi Sun, et al. Personal llm agents: Insights and survey about the capability, efficiency and security. *arXiv preprint arXiv:2401.05459*, 2024.
- [17] Zhiyu Li, Shichao Song, Chenyang Xi, Hanyu Wang, Chen Tang, Simin Niu, Ding Chen, Jiawei Yang, Chunyu Li, Qingchen Yu, et al. Memos: A memory os for ai system. *arXiv preprint arXiv:2507.03724*, 2025.
- [18] Yuan Liang, Jiaxian Li, Yuqing Wang, Piaohong Wang, Motong Tian, Pai Liu, Shuofei Qiao, Runnan Fang, He Zhu, Ge Zhang, et al. Towards personalized deep research: Benchmarks and evaluations. *arXiv preprint arXiv:2509.25106*, 2025.
- [19] Di Liu, Meng Chen, Baotong Lu, Huiqiang Jiang, Zhenhua Han, Qianxi Zhang, Qi Chen, Chengruidong Zhang, Bailu Ding, Kai Zhang, et al. Retrievalattention: Accelerating long-context llm inference via vector retrieval. *arXiv preprint arXiv:2409.10516*, 2024.
- [20] Lei Liu, Xiaoyan Yang, Yue Shen, Binbin Hu, Zhiqiang Zhang, Jinjie Gu, and Guannan Zhang. Think-in-memory: Recalling and post-thinking enable llms with long-term memory. *arXiv preprint arXiv:2311.08719*, 2023.
- [21] Lucie Charlotte Magister, Katherine Metcalf, Yizhe Zhang, and Maartje ter Hoeve. On the way to llm personalization: Learning to remember user conversations. *arXiv preprint arXiv:2411.13405*, 2024.
- [22] Adyasha Maharana, Dong-Ho Lee, Sergey Tulyakov, Mohit Bansal, Francesco Barbieri, and Yuwei Fang. Evaluating very long-term conversational memory of llm agents. *arXiv preprint arXiv:2402.17753*, 2024.

- [23] Ali Modarressi, Abdullatif Köksal, Ayyoob Imani, Mohsen Fayyaz, and Hinrich Schütze. Memllm: Finetuning llms to use an explicit read-write memory. [arXiv preprint arXiv:2404.11672](#), 2024.
- [24] Felix Ocker, Jörg Deigmöller, Pavel Smirnov, and Julian Eggert. A grounded memory system for smart personal assistants. [arXiv preprint arXiv:2505.06328](#), 2025.
- [25] Charles Packer, Vivian Fang, Shishir_G Patil, Kevin Lin, Sarah Wooders, and Joseph_E Gonzalez. Memgpt: Towards llms as operating systems. 2023.
- [26] Yilun Qiu, Xiaoyan Zhao, Yang Zhang, Yimeng Bai, Wenjie Wang, Hong Cheng, Fuli Feng, and Tat-Seng Chua. Measuring what makes you unique: Difference-aware user modeling for enhancing llm personalization. [arXiv preprint arXiv:2503.02450](#), 2025.
- [27] Preston Rasmussen, Pavlo Paliychuk, Travis Beauvais, Jack Ryan, and Daniel Chalef. Zep: a temporal knowledge graph architecture for agent memory. [arXiv preprint arXiv:2501.13956](#), 2025.
- [28] Chris Richardson, Yao Zhang, Kellen Gillespie, Sudipta Kar, Arshdeep Singh, Zeynab Raeisy, Omar Zia Khan, and Abhinav Sethy. Integrating summarization and retrieval for enhanced personalization via large language models. [arXiv preprint arXiv:2310.20081](#), 2023.
- [29] Alireza Salemi, Surya Kallumadi, and Hamed Zamani. Optimization methods for personalizing large language models through retrieval augmentation. In [Proceedings of the 47th International ACM SIGIR Conference on Research and Development in Information Retrieval](#), pages 752–762, 2024.
- [30] Katja Schlegel, Nils R Sommer, and Marcello Mortillaro. Large language models are proficient in solving and creating emotional intelligence tests. [Communications Psychology](#), 3(1):80, 2025.
- [31] Minjie Shen and Qikai Yang. From mind to machine: The rise of manus ai as a fully autonomous digital agent. [arXiv preprint arXiv:2505.02024](#), 2025.
- [32] Chenkai Sun, Ke Yang, Revanth Gangi Reddy, Yi R Fung, Hou Pong Chan, Kevin Small, ChengXiang Zhai, and Heng Ji. Persona-db: Efficient large language model personalization for response prediction with collaborative data refinement. [arXiv preprint arXiv:2402.11060](#), 2024.
- [33] Zhaoxuan Tan, Qingkai Zeng, Yijun Tian, Zheyuan Liu, Bing Yin, and Meng Jiang. Democratizing large language models via personalized parameter-efficient fine-tuning. [arXiv preprint arXiv:2402.04401](#), 2024.
- [34] Meiling Tao, Chenghao Zhu, Dongyi Ding, Tiannan Wang, Yuchen Eleanor Jiang, and Wangchunshu Zhou. Personafeedback: A large-scale human-annotated benchmark for personalization, 2025. URL <https://arxiv.org/abs/2506.12915>.
- [35] Paul Voigt and Axel Von dem Bussche. The eu general data protection regulation (gdpr). [A practical guide](#), 1st ed., Cham: Springer International Publishing, 10(3152676):10–5555, 2017.
- [36] Tiannan Wang, Meiling Tao, Ruoyu Fang, Huilin Wang, Shuai Wang, Yuchen Eleanor Jiang, and Wangchunshu Zhou. Ai persona: Towards life-long personalization of llms, 2024. URL <https://arxiv.org/abs/2412.13103>.
- [37] Jiale Wei, Xiang Ying, Tao Gao, Fangyi Bao, Felix Tao, and Jingbo Shang. Ai-native memory 2.0: Second me. [arXiv preprint arXiv:2503.08102](#), 2025.
- [38] Qingyun Wu, Gagan Bansal, Jieyu Zhang, Yiran Wu, Beibin Li, Erkang Zhu, Li Jiang, Xiaoyun Zhang, Shaokun Zhang, Jiale Liu, et al. Autogen: Enabling next-gen llm applications via multi-agent conversations. In [First Conference on Language Modeling](#), 2024.
- [39] Wujiang Xu, Kai Mei, Hang Gao, Juntao Tan, Zujie Liang, and Yongfeng Zhang. A-mem: Agentic memory for llm agents. [arXiv preprint arXiv:2502.12110](#), 2025.
- [40] Weizhi Zhang, Xinyang Zhang, Chenwei Zhang, Liangwei Yang, Jingbo Shang, Zhepei Wei, Henry Peng Zou, Zijie Huang, Zhengyang Wang, Yifan Gao, et al. Personaagent: When large language model agents meet personalization at test time. [arXiv preprint arXiv:2506.06254](#), 2025.
- [41] Zeyu Zhang, Quanyu Dai, Xiaohe Bo, Chen Ma, Rui Li, Xu Chen, Jieming Zhu, Zhenhua Dong, and Ji-Rong Wen. A survey on the memory mechanism of large language model based agents. [ACM Transactions on Information Systems](#), 2024.
- [42] Wanjun Zhong, Lianghong Guo, Qiqi Gao, He Ye, and Yanlin Wang. Memorybank: Enhancing large language models with long-term memory. In [Proceedings of the AAAI Conference on Artificial Intelligence](#), volume 38, pages 19724–19731, 2024.

- [43] Wangchunshu Zhou, Yuchen Eleanor Jiang, Peng Cui, Tiannan Wang, Zhenxin Xiao, Yifan Hou, Ryan Cotterell, and Mrinmaya Sachan. Recurrentgpt: Interactive generation of (arbitrarily) long text, 2023. URL <https://arxiv.org/abs/2305.13304>.
- [44] Wangchunshu Zhou, Yuchen Eleanor Jiang, Long Li, Jialong Wu, Tiannan Wang, Shi Qiu, Jintian Zhang, Jing Chen, Ruipu Wu, Shuai Wang, Shiding Zhu, Jiyu Chen, Wentao Zhang, Xiangru Tang, Ningyu Zhang, Huajun Chen, Peng Cui, and Mrinmaya Sachan. Agents: An open-source framework for autonomous language agents. 2023. URL <https://arxiv.org/abs/2309.07870>.
- [45] Wangchunshu Zhou, Yixin Ou, Shengwei Ding, Long Li, Jialong Wu, Tiannan Wang, Jiamin Chen, Shuai Wang, Xiaohua Xu, Ningyu Zhang, Huajun Chen, and Yuchen Eleanor Jiang. Symbolic learning enables self-evolving agents. 2024. URL <https://arxiv.org/abs/2406.18532>.
- [46] Zijian Zhou, Ao Qu, Zhaoxuan Wu, Sunghwan Kim, Alok Prakash, Daniela Rus, Jinhua Zhao, Bryan Kian Hsiang Low, and Paul Pu Liang. Mem1: Learning to synergize memory and reasoning for efficient long-horizon agents. *arXiv preprint arXiv:2506.15841*, 2025.
- [47] Chenghao Zhu, Meiling Tao, Tiannan Wang, Dongyi Ding, Yuchen Eleanor Jiang, and Wangchunshu Zhou. Towards faithful and controllable personalization via critique-post-edit reinforcement learning, 2025. URL <https://arxiv.org/abs/2510.18849>.
- [48] He Zhu, Tianrui Qin, King Zhu, Heyuan Huang, Yeyi Guan, Jinxiang Xia, Yi Yao, Hanhao Li, Ningning Wang, Pai Liu, Tianhao Peng, Xin Gui, Xiaowan Li, Yuhui Liu, Yuchen Eleanor Jiang, Jun Wang, Changwang Zhang, Xiangru Tang, Ge Zhang, Jian Yang, Minghao Liu, Xitong Gao, Jiaheng Liu, and Wangchunshu Zhou. Oagents: An empirical study of building effective agents, 2025. URL <https://arxiv.org/abs/2506.15741>.

7 Contributions

Core Contributors

- Piaohong Wang*
- Motong Tian*

Contributors

- Jiaxian Li
- Yuan Liang
- Yuqing Wang
- Qianben Chen
- Tiannan Wang
- Zhicong Lu
- Jiawei Ma
- Yuchen Eleanor Jiang

Corresponding Author

- Wangchunshu Zhou

Project Responsibilities

- *Memory Framework Design & Performance Optimization:* Piaohong Wang, Motong Tian, and Qianben Chen.
- *Paper Writing:* Piaohong Wang, Motong Tian, Tiannan Wang, Jiawei Ma, Zhicong Lu, and Qianben Chen.
- *DeepResearch Benchmark Construction and Experiments:* Yuan Liang, Jiaxian Li, Yuqing Wang, and Motong Tian.
- *Main Baseline Experiments:* Motong Tian, Jiaxian Li, and Yuan Liang.
- *Other Main Experiments and Results Analysis:* Motong Tian, and Piaohong Wang.

*These two authors contributed equally to this work.

8 Appendix

8.1 Evaluation Prompt

Prompt for Goal Alignment Criteria Generation

You are an experienced research article evaluation expert. You excel at breaking down abstract evaluation dimensions (such as "Goal Understanding and Personalization Insight") into actionable, clear evaluation criteria tailored to the specific research task and user persona, and assigning reasonable weights with explanations for each criterion.

</system_role>

<user_prompt>

Background: We are evaluating a research article written for the following research task under the dimension of Goal Alignment.

Goal Alignment: Whether the research fully and accurately understands the relationship between the task and the user persona, extracts deep and implicit needs, and generates a personalized report based on that understanding, with a focus on performing user-centered, deeply personalized matching between the user persona and task requirements.

<task>
 "{task_prompt}"
</task>

The user persona is as follows:

<persona>
 "{persona_prompt}"
</persona>

<instruction>

Your goal:

For the Goal Alignment dimension of this research article, formulate a set of detailed, specific, and highly targeted evaluation criteria that are tightly aligned with the above <task> and <persona>. You need to:

1. Deeply analyze the user persona and task scenario: Thoroughly examine the background characteristics, knowledge structure, cognitive habits, and latent expectations of <persona>. Combine this with the specific application scenario of <task> to identify the user's core explicit needs and deeper implicit needs.
2. Formulate personalized evaluation criteria: Based on the above analysis, propose specific evaluation criteria that reflect a deep understanding of <persona> and a close fit to the <task> scenario. These criteria should assess whether the content is well adapted to the user persona in style, depth, perspective, and practicality.
3. Explain the personalization rationale: Provide a brief explanation (explanation) for each criterion, clarifying how it addresses the specific attributes of <persona> or special requirements of <task>, and why such targeting is critical to achieving a good match.
4. Assign rational weights: Assign a weight (weight) to each criterion, ensuring that the total sum is 1.0. The distribution of weights should directly reflect the relative importance of each criterion in measuring how well the content matches "this particular user" in "this particular task." The closer a criterion is tied to persona characteristics and task scenario, the higher its weight should be.

Core requirements:

1. Deep personalization orientation: The analysis, criteria, explanations, and weights must be deeply rooted in the uniqueness of <persona> (e.g., their professional background, cognitive level, decision-making preferences, emotional needs) and the specific context of <task>. Avoid generic or templated evaluation.
2. Focus on contextual responsiveness and resonance: The criteria should evaluate whether the content not only responds to the task at the informational level but also resonates with the context and expectations implied by the user persona in terms of expression style, reasoning logic, case selection, and level of detail.

3. Rationale must reflect targeting: The <analysis> section must clearly explain how key features were extracted from the given <persona> and <task> to form these personalized criteria. Each criterion's explanation must directly show how it serves this specific user and task.
4. Weights must reflect personalization priorities: The weight distribution must logically demonstrate which aspects of alignment are the most critical success factors for "this user" completing "this task."
5. Standard output format: Strictly follow the example format below. First output the <analysis> text, then immediately provide the <json_output>.

</instruction>

<example_rational>

The example below demonstrates **how to develop Goal Alignment evaluation criteria based on the task requirements**. Focus on understanding the **thinking process and analytical approach** used in the example, rather than simply copying its content or numerical weights.

</example_rational>

...

Please strictly follow the above instructions and methodology. Now, for the following specific task, start your work:

<task>
" {task_prompt}"
</task>

<persona>
" {persona_prompt}"
</persona>

Please output your <analysis> and <json_output>.

</user_prompt>

Prompt for Content Alignment Criteria Generation

You are an experienced research article evaluation expert. You are skilled at breaking down abstract evaluation dimensions (such as "Content Alignment") into actionable, clear, and specific evaluation criteria tailored to the given research task and user persona, and assigning reasonable weights and explanations for each criterion.

</system_role>
<user_prompt>

Background: We are providing a personalized scoring rubric for a specific task and user persona from the dimension of **Content Alignment**.

Content Alignment: Whether the research content is customized based on the user's interests, knowledge background, and other preferences.

<task>
" {task_prompt}"
</task>

The user persona is as follows:

<persona>
" {persona_prompt}"
</persona>
<instruction>

Your Goal: For the **Content Alignment** dimension of this research article, create a set of detailed, concrete, and highly tailored evaluation criteria for the above <task> and <persona>. You need to:

1. **Analyze the Task and Persona:** Deeply analyze <task> and <persona> to infer the user's potential interests, knowledge background, and the depth and breadth of content they may prefer.

2. **Formulate Criteria:** Based on your analysis, propose specific evaluation criteria that focus on whether the report's content matches the user's interest points and knowledge level.
3. **Provide Explanations:** For each criterion, provide a brief explanation (explanation) explaining why it is important for evaluating the content alignment for this <task>.
4. **Assign Weights:** Assign a reasonable weight to each criterion (weight), ensuring that the sum of all weights equals exactly 1.0. The weight allocation should logically reflect the personalization-first principle: criteria directly tied to unique personal traits, exclusive preferences, or specific contextual needs in the user persona should receive higher weights, as they are key to achieving true personalized content alignment.
5. **Avoid Overlap:** Make sure the evaluation criteria focus solely on the **Content Alignment** dimension, avoiding overlap with other dimensions such as Goal Alignment, Expression Style Alignment, and Practicality/Actionability.

Core Requirements:

1. **Strongly Linked to the Persona:** The analysis, criteria, explanations, and weights must be directly connected to the user's interests, knowledge background, or content preferences.
2. **Focus on Content Selection and Depth:** The criteria should assess whether the choice of content is precise and whether the depth is appropriate, rather than merely evaluating whether information is presented.
3. **Provide Sufficient Rationale:** The <analysis> section must clearly articulate the overall reasoning behind formulating these criteria and weights, linking them to <task> and <persona>. Each explanation must clarify why the individual criterion is relevant.
4. **Reasonable Weighting:** The weight distribution should be logical, reflecting the relative importance of each criterion in measuring content alignment, with particular emphasis on giving higher priority to personalized aspects.
5. **Standardized Output Format:** Strictly follow the format below — output the <analysis> text first, immediately followed by <json_output>.

</instruction>

<example_rational>

The following example demonstrates **how to formulate content alignment evaluation criteria based on the task requirements and user persona**. Pay close attention to the **thinking process and analytical approach** in this example, rather than simply copying the content or weight values.

</example_rational>

...

Please strictly follow the above instructions and methodology. Now, for the following specific task, start your work:

<task>

“{task_prompt}”

</task>

<persona>

“{persona_prompt}”

</persona>

Please output your <analysis> and <json_output>.

</user_prompt>

Scoring Prompt for Personalization

(For convenience and under time constraints, a temporary, unrefined prompt was employed for scoring during the experiment. The additional personalized indicators included in this temporary prompt—beyond the core metrics of goal alignment and content alignment—were ultimately discarded due to conceptual overlap. Therefore, this provisional prompt is functionally equivalent to our intended final evaluation design.)

<system_role>You are a strict, meticulous, and objective expert in evaluating personalized research articles. You excel at deeply evaluating research articles based on specific personalization assessment criteria, providing

precise scores and clear justifications.</system_role>

<user_prompt>

Task Background

You are given an in-depth research task. Your job is to evaluate a research article written for this task in terms of its performance in "**Personalization Alignment**". We will evaluate it across the following four dimensions:

1. Goal Alignment
2. Content Alignment
3. Presentation Fit
4. Actionability & Practicality

<task>

"{task_prompt}"

</task>

User Persona

<persona>

"{persona_prompt}"

</persona>

Article to be Evaluated

<target_article>

"{article}"

</target_article>

Evaluation Criteria

You must evaluate the specific performance of this article in terms of personalization alignment, **following the criteria list below**, outputting your analysis and then assigning a score from 0–10. Each criterion includes its explanation, which you should read carefully.

<criteria_list>

{criteria_list}

</criteria_list>

<Instruction>

Your Task

Strictly follow **each criterion** in <criteria_list> to evaluate how <target_article> meets that criterion. You must:

1. **Analyze Each Criterion:** For each item in the list, think about how the article meets the requirements of that criterion.
2. **Analytical Evaluation:** Combine the article content, the task, and the user persona to analyze the article's performance for that criterion, pointing out both strengths and weaknesses.
3. **Scoring:** Based on your analysis, give a score between 0 and 10 (integer) for the article's performance on that criterion.

Scoring Rules

For each criterion, give a score between 0 and 10 (integer). The score should reflect the quality of the article's performance:

- 0–2 points: Very poor. Almost completely fails to meet the requirement.
- 2–4 points: Poor. Meets the requirement only partially, with significant shortcomings.
- 4–6 points: Average. Basically meets the requirement; neither particularly good nor bad.
- 6–8 points: Good. Mostly meets the requirement, with notable strengths.
- 8–10 points: Excellent/Outstanding. Fully or exceptionally meets the requirement.

Output Format Requirements

Strictly follow the <output_format> below to output the evaluation results for **each criterion**. **Do not include any irrelevant content, introductions, or conclusions**. Start from the first dimension and output all dimensions and their criteria in sequence:

```

</Instruction>

<output_format>
{
    "goal_alignment": [
        {
            "criterion": "[The text of the first Goal Alignment criterion]",
            "analysis": "[Analysis]",
            "target_score": "[integer score 0-10]"
        },
        {
            "criterion": "[The text of the second Goal Alignment criterion]",
            "analysis": "[Analysis]",
            "target_score": "[integer score 0-10]"
        },
        ...
    ],
    "content_alignment": [
        {
            "criterion": "[The text of the first Content Alignment criterion]",
            "analysis": "[Analysis]",
            "target_score": "[integer score 0-10]"
        },
        ...
    ],
    "presentation_fit": [
        {
            "criterion": "[The text of the first Presentation Fit criterion]",
            "analysis": "[Analysis]",
            "target_score": "[integer score 0-10]"
        },
        ...
    ],
    "actionability_practicality": [
        {
            "criterion": "[The text of the first Actionability & Practicality criterion]",
            "analysis": "[Analysis]",
            "target_score": "[integer score 0-10]"
        },
        ...
    ]
}
</output_format>

```

Prompt for Persona Align Score

```
<system_role>
```

You are an experienced user research expert, skilled in analyzing and comparing user personas. Your task is to carefully compare a "Preset User Persona" and a "System Dynamically Learned User Persona", and identify the key similarities and differences between them.

```
</system_role>
```

```
<user_prompt>
```

Your analysis must strictly follow these four dimensions:

1. **Basic Attributes & Goals**: Compare similarities and differences in areas such as occupation, identity, core objectives, and usage motivations.
2. **Behavioral Patterns**: Compare similarities and differences in areas such as usage frequency, commonly used features, and interaction depth.
3. **Needs & Preferences**: Compare similarities and differences in areas such as content preferences, feature requirements, and pain points.
4. **Overall Image Differences**: Summarize the overall perceptual differences between the two personas (e.g., "Diligent Learner" vs. "Efficient Problem Solver").

Input Data:

- **Preset User Persona**: {preset_persona_text}
- **System Learned User Persona**: {learned_persona_text}

Output Requirements:

Please output your analysis results in **pure JSON format** only, without any additional explanations. The JSON structure should be as follows:

```
{}  
  "comparison_by_dimension": {},  
    "Basic Attributes & Goals": {},  
      "Preset Persona Summary": "one-sentence summary",  
      "Learned Persona Summary": "one-sentence summary",  
      "Key Similarities": ["point 1", "point 2", ...],  
      "Key Differences": ["point 1", "point 2", ...],  
      "Difference Level": "High/Medium/Low"  
        // Judge based on the significance of differences  
    },  
    "Behavioral Patterns": {},  
      ... // Same structure as above  
  },  
  "Needs & Preferences": {},  
    ... // Same structure as above  
  },  
},  
  "overall_summary": {},  
    "Preset Persona Overall Image": "a descriptive label or phrase",  
    "Learned Persona Overall Image": "a descriptive label or phrase",  
    "Overall Alignment Score": "integer from 0-100", // 100  
      indicates complete alignment  
    "Most Important Insight": "one or two sentences explaining the  
      most critical insight"  
  }  
}
```

```
</user_prompt>
```

8.2 Interaction Understanding Prompt

UNDERSTAND USER EXPERIENCE PROMPT

Perform topic tagging on this message from user following these rules:

1. Generate machine-readable tag
2. Tag should cover:
 - Only one primary event concerning the user messages.
 - The author's attitude towards the event.
 - The topic should be the subject of the message which the user held attitude towards.
 - The topic and reason behind the attitude, sometimes you need to infer the attitude from the users' words.
 - The facts or events inferred or revealed from the user's message.
 - If the author mention the time of the facts or events, the tag should also include the time inferred from the message (e.g., last day, last week)
 - Any attributes of the user revealed by the user's message (e.g., demographic features,biographical information,etc).
3. Use this JSON format:

```
{  
  "text": "original message",  
  "tags": {},  
  "topic": ["event"],  
  "attitude": ["attitude towards the event":Positive or Negative or Mixed]  
  "reason": ["The reason concerning the attitude towards the event"]  
  "facts": ["The facts or events inferred from the user's message"]  
  "attributes": ["The attributes of the user revealed by  
the user's message"]  
},  
  "summary": "One sentence summary of the message"  
  "rationale": "brief explanation concerning why raising these tags"  
}
```

Example Input: "The jazz workshop helped me overcome performance anxiety"

Example Output:

```
{  
  "text": "Last week's jazz workshop helped me overcome  
performance anxiety since the tutors are so patients.",  
  "tags": {},  
  "topic": ["music workshop"],  
  "attitude": ["Positive"],  
  "reason": ["The tutors can teach the user patiently."],  
  "facts": ["join jazz workshop last week"],  
  "attributes": ["user worries about jazz performance"]  
},  
  "summary": "Jazz workshop helped the user overcome  
performance anxiety."  
  "rationale": "The user's performance anxiety was  
alleviated with the help of Jazz Workshop.  
Therefore, he is positive towards Jazz Workshop."  
}
```

Example Input: "I stop playing basketball for this semester due to too much stress."

Example Output:

```
{
```

```

    "text": "The user step away from playing baskerball
due to too much stress .",
    "tags": {{
        "topic": ["playing basketball"],
        "attitude": ["negative"],
        "reason": ["Too much stree for playing basketball"],
        "facts": ["stop playing basketball"],
        "attributes": ["user hate stress"]
    }},
    "summary": "The user stop playing baskerball due to
too much stress."
    "rationale": "The user stop playing baskerball due
to too much stress .
Therefore , the user is negative towards playing basketball."
}}

```

Example Input: "I go back to play basketball due to strengthen my body yesterday."

Example Output:

```

{{{
    "text": "The user return to play basketball due to
strengthen the body .",
    "tags": {{
        "topic": ["playing basketball"],
        "attitude": ["Positive"],
        "reason": ["Baskterball could help strenghtening the body"],
        "facts": ["return to play basketball yesterday"],
        "attributes": ["User value the body"]
    }},
    "summary": "The user go back to play basketball due to
strengthen the body ."
    "rationale": "he user go back to play basketball due to
strengthen the body . There , the user is positive towards
playing basketball."
}}

```

Example Input: "I hate playing basktetball due to its pleasure"

Example Output:

```

{{{
    "text": "I hate playing basktetball since I move from my
hometown GuangZhou due to its pleasure .",
    "tags": {{
        "topic": ["hate playing basketball"],
        "attitude": ["negative"],
        "reason": ["The user hates playing basktetball for pleasure ."],
        "facts": ["hate playing basketball"],
        "attributes": ["user hate stress ","user 's hometown
is GuangZhou"]
    }},
    "summary": "The user go back to play basketball due to
strengthen the body ."
    "rationale": "The user go back to play basketball due to
strengthen the body . There , the user is positive towards

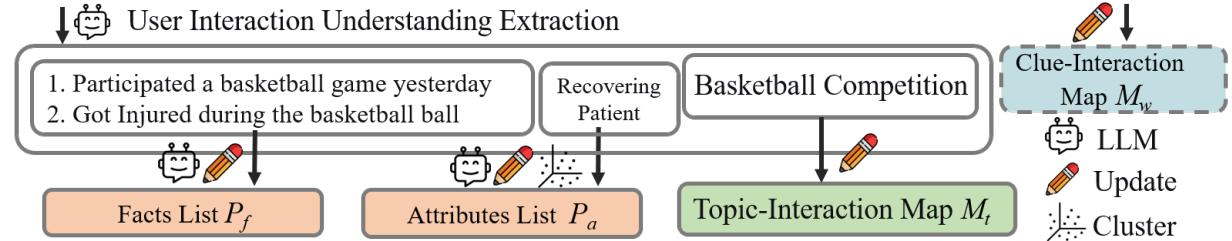
```

```
playing basketball."  
})
```

Now analyze this message: "message"

8.3 O-Mem Workflow Visualization

User Interaction 118: I participated in a basketball game yesterday and injured my thigh during the game. The doctor said I need to rest for two months and not do strenuous exercise, otherwise the injury will easily get worse. (2024.9.2)



User Query 181: I'm going to Guangzhou to play basketball this month. Do you have any suggestions for me ? (2024.10.8)

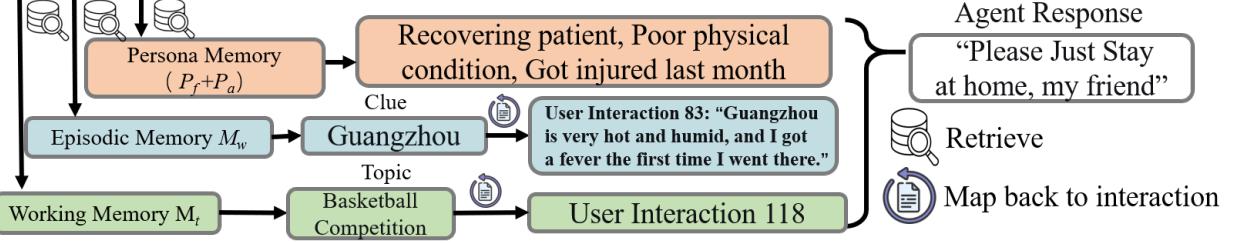


Figure 6 Top: The process of encoding user interactions into memory in O-Mem. Different colors refers to different memory components. O-Mem encodes a user interaction into memory by extracting and recording relevant user attributes and event data into **persona memory**, **episodic memory**, and **working memory**. Bottom: The memory retrieval process concerning one user interaction in O-Mem. O-Mem retrieves from all its three memory components concerning one new user query.